

# CAS in Applied Data Science Muerren 2024



Géraldine Schaller-Conti

# Program of the week

17:00 - 18:30 Tutorial (Geraldine) **Convolutions**

19:00 - 20:00 Dinner at hotel Regina

20:30 - 21:00 Neural Style Transformations (an example on how to use deep ML to create art - Sigve)

## Thursday (Geraldine)

08:15 - 09:00 Lecture **Deep Forward Networks – optimization**

09:15 - 10:00 Tutorial **Transfer Learning**

10:00 - 10:30 Coffee break

10:30 - 12:30 Tutorial **Transfer Learning**

12:30 - 17:00 Skiing, work or whatever

17:00 - 18:15 Tutorial and project discussions **Projects**

19:00 - 22:00 Fondue and sledge ride back down to Regina in the dark

Everyone has to be READY OUTSIDE THE HOTEL at 18:45.

The train up departs at 19:00.

## Friday (Geraldine)

08:15 - 09:00 Lecture **Convolutional models and generative models**

09:15 - 10:00 Tutorial **Recurrent Neural Networks**

10:00 - 10:30 Coffee break and Check Out

10:30 - 12:00 Tutorial / Discussion Session **Recurrent Neural Networks**

12:00 - 12:30 Wrap up (Sigve)

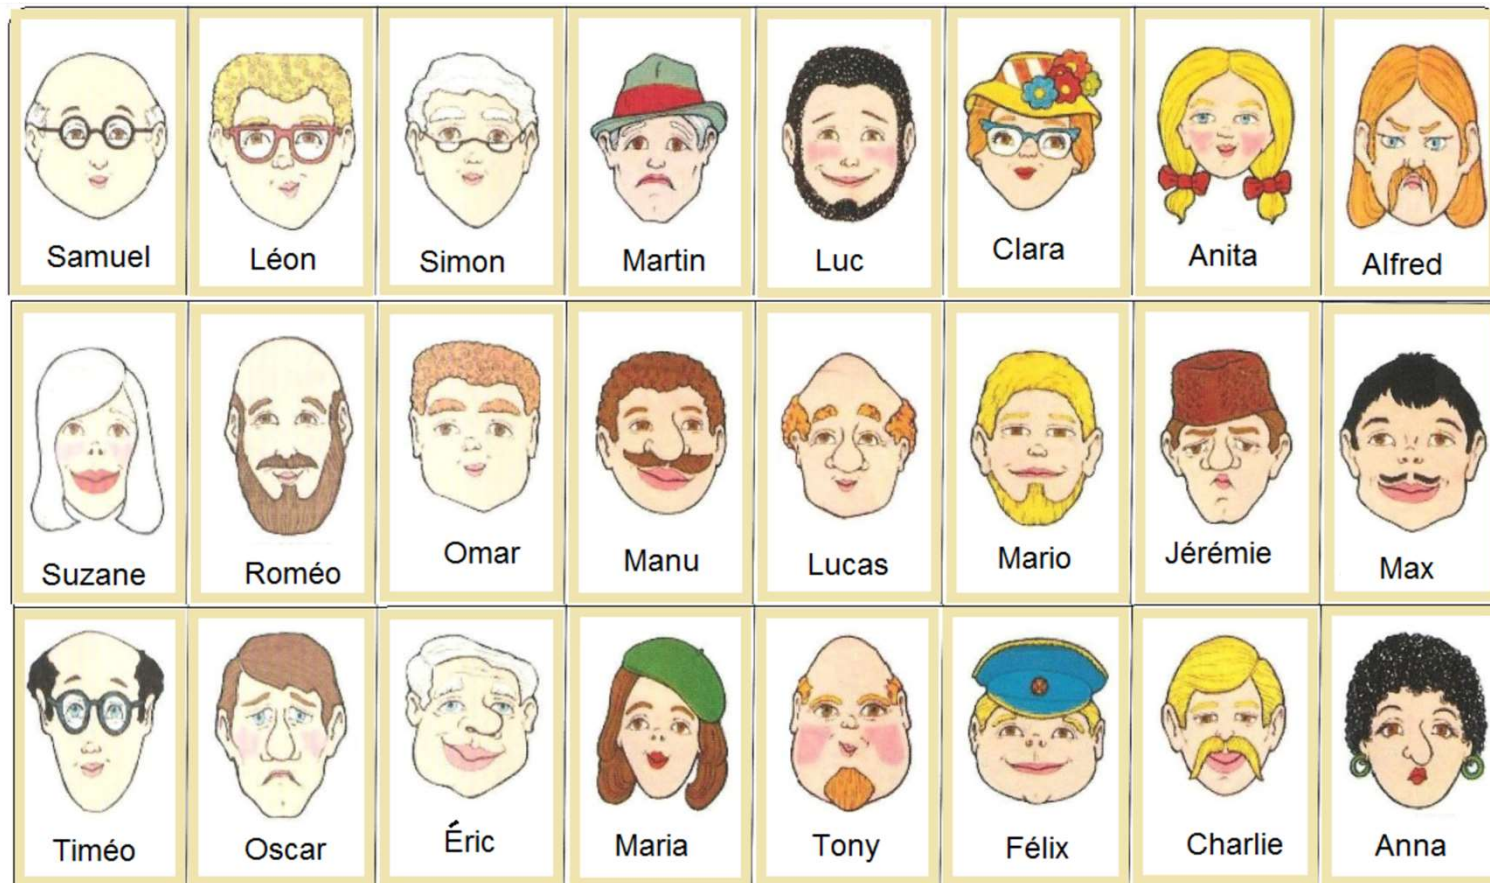
12:30 End

# About me ☺



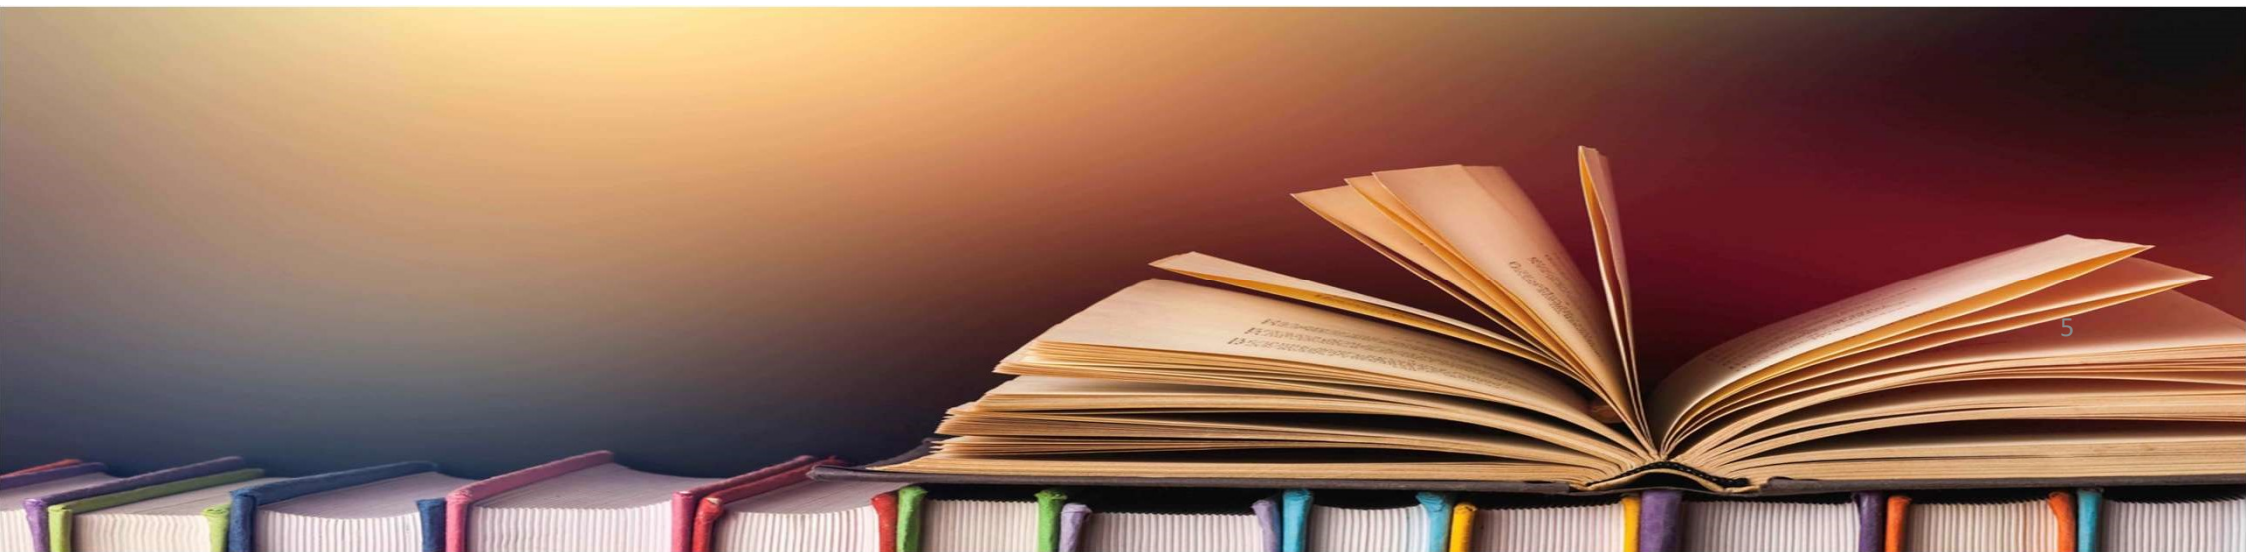


# About you ☺



# Bibliography

- Deep Learning book (Goodfellow, Bengio, Courville)
- Machine Learning @ Stanford (Prof Andrew Ng)
- Hands-On Machine Learning with Scikit-Learn & Tensorflow (Aurélien Géron)



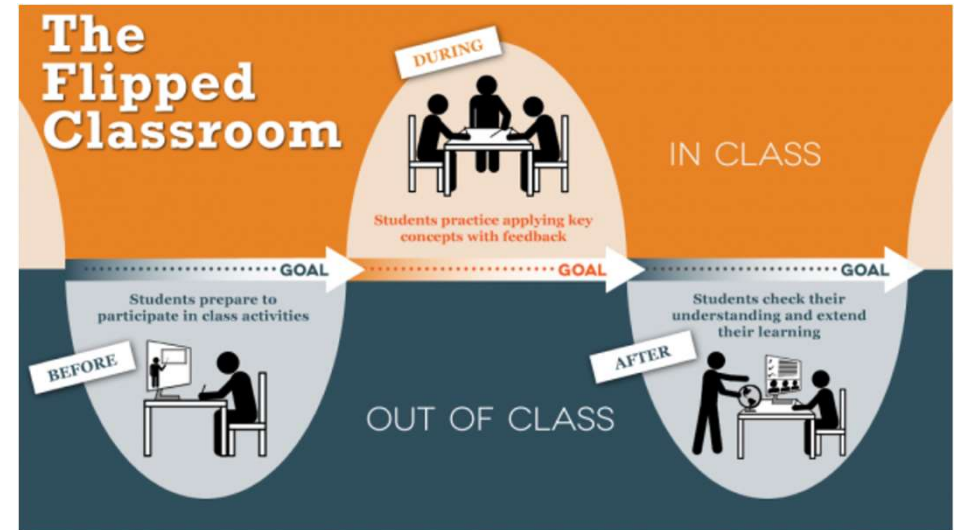
# Teaching method

## Inverted classroom based

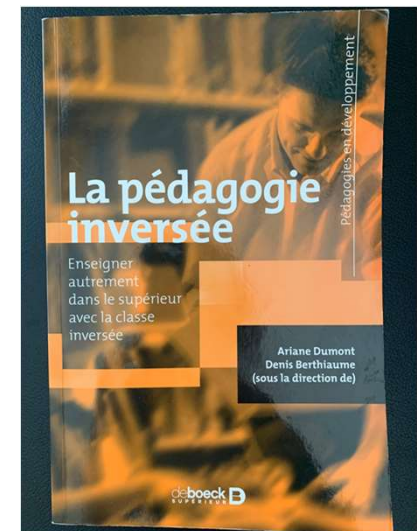
- Introduction lectures
- Real content you learn yourself with the **notebooks**. *Either to put in practice your knowledge or to learn ahead of another lecture*

## Why

- Supposed to be better
- More fun
- Learning by doing



*To give back sense to being present (Marcel Lebrun)*



# Tutorial IV : Convolutions

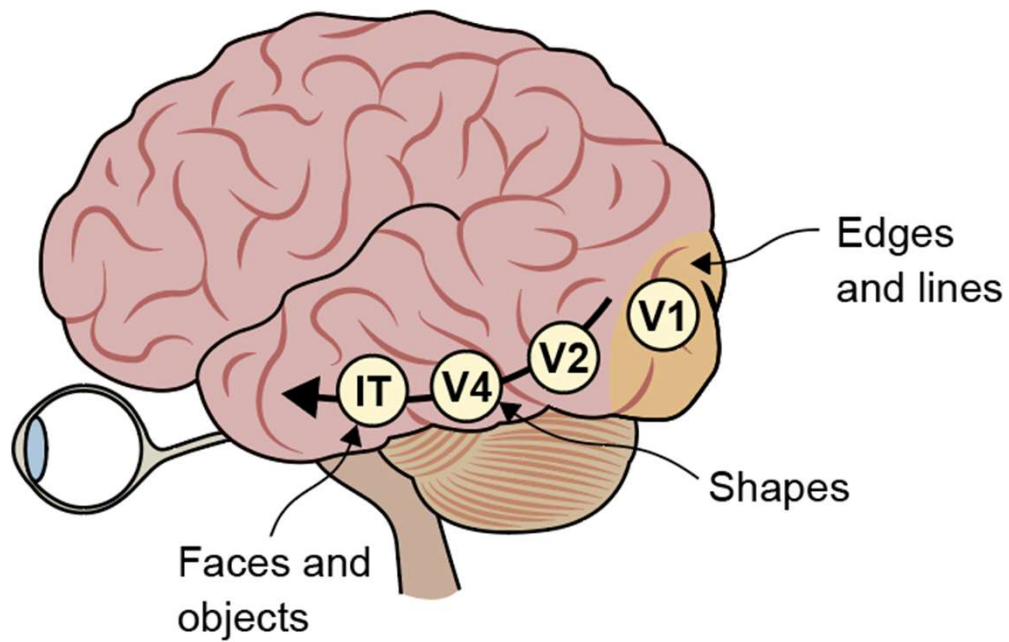
# Introduction

- **Goal** : basics to perform image recognition
- **Program** : inverted classroom style
  - Theory
  - Overview to get the big picture of the Notebook
  - Work alone (30 minutes)
  - Work in groups of 4 ( 30 minutes)
- **Technical** : Google Colab, Pytorch

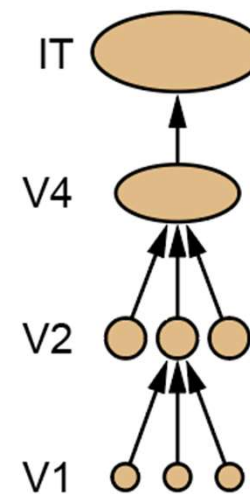


# Theory


# Human vision



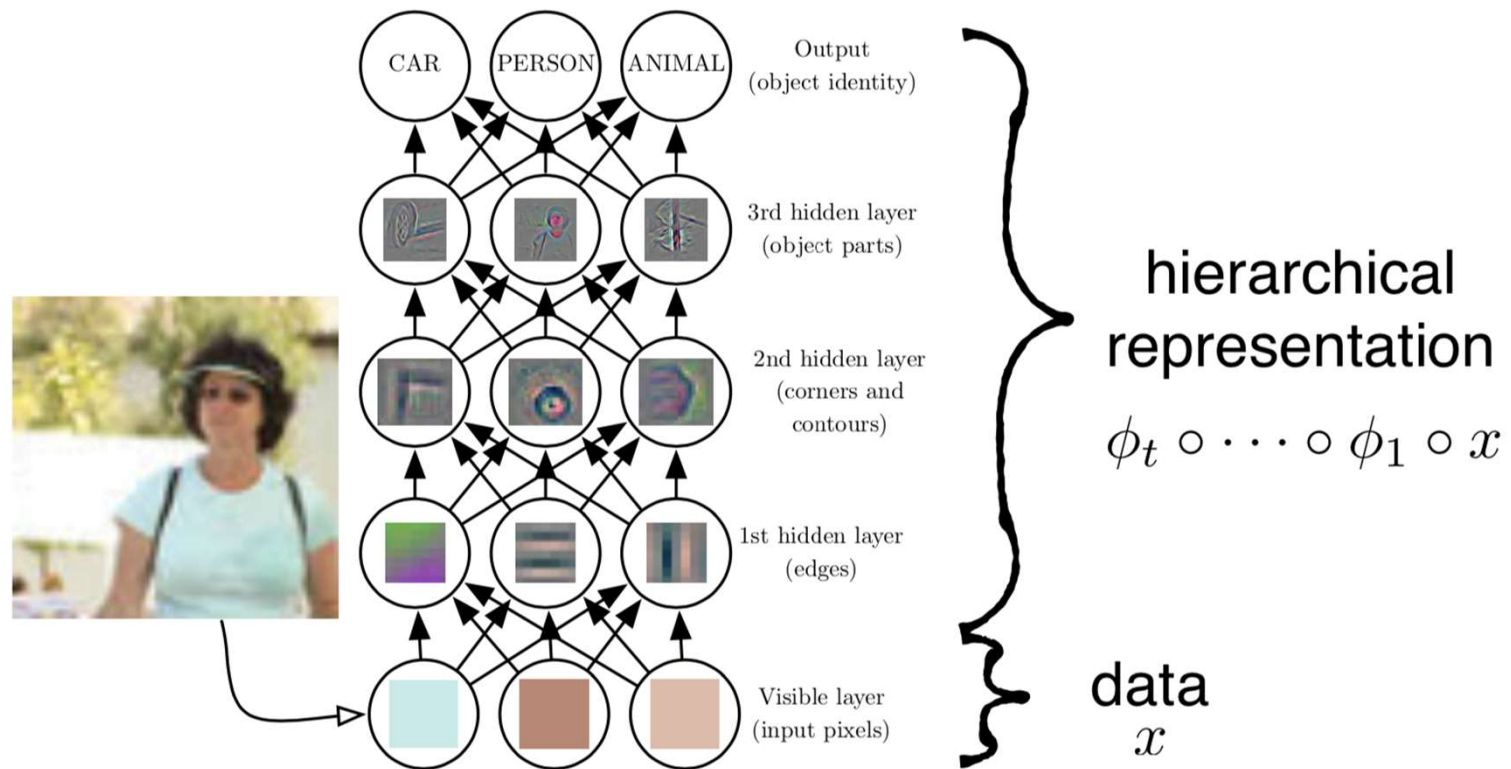
Receptive fields size



Features



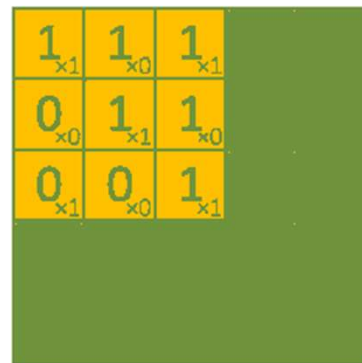
# Computer vision



# Kernel (filter)

- Used to **detect features** (vertical/horizontal filter,...)
- Different kernels to create different feature maps → learn to see various patterns and details in images

1	0	1
0	1	0
1	0	1



Image

4		

Convolved  
Feature

how well the pattern in the kernel matches  
the content in that part of the image

= Feature map


# Kernel example

1	0	-1
1	0	-1
1	0	-1

Vertical

1	1	1
0	0	0
-1	-1	-1

Horizontal



Vertical edges



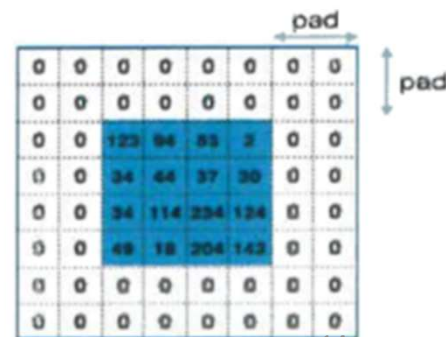
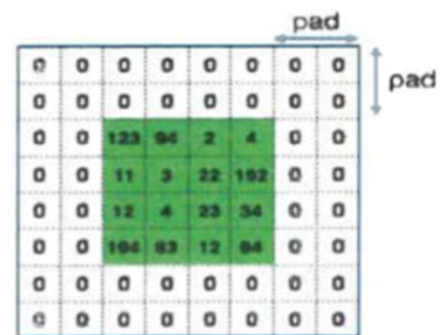
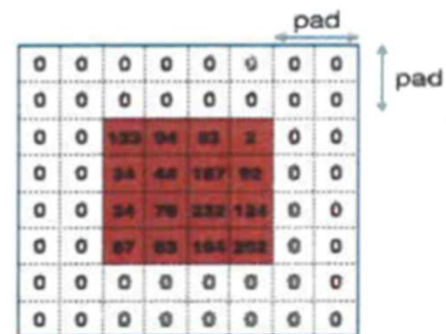
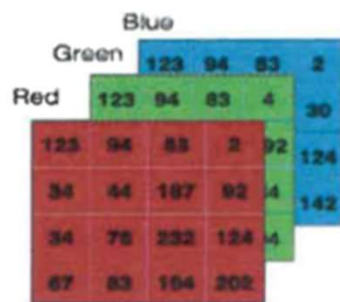
Horizontal edges



# Padding, Stride, Dilation



=



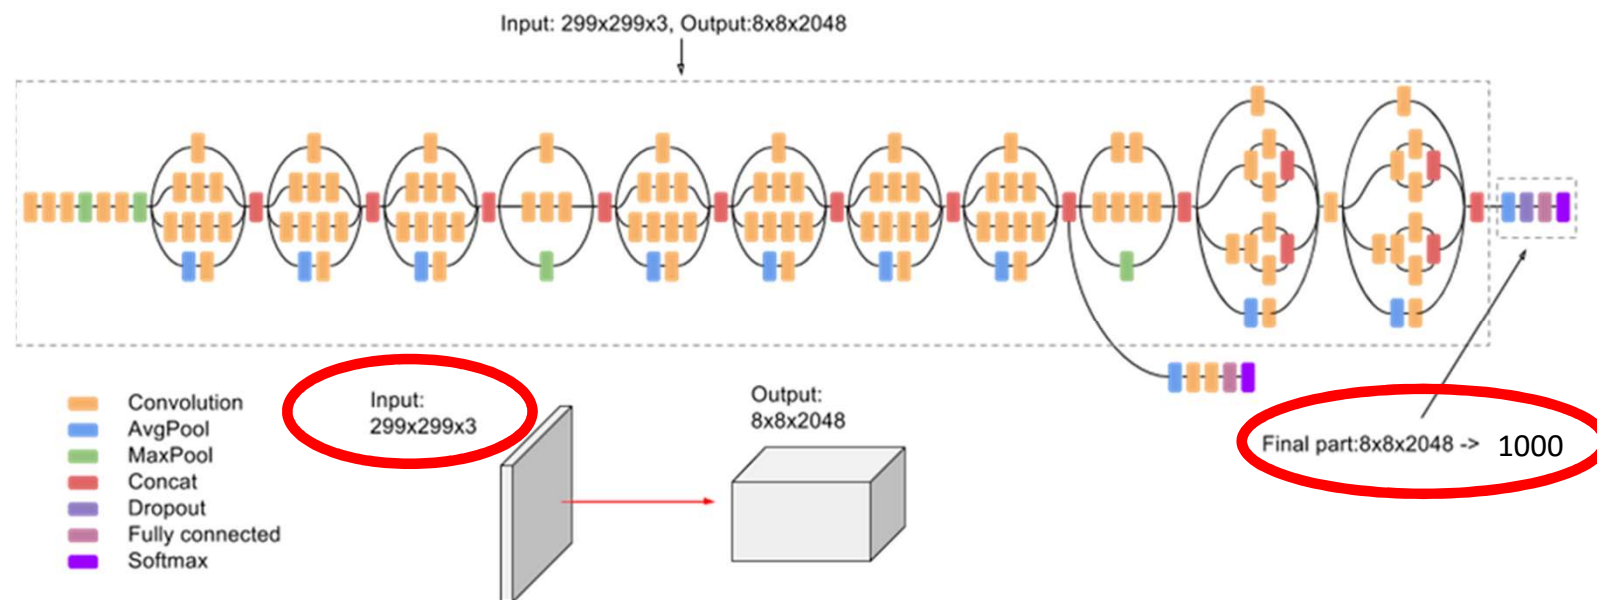
1	2	3
4	5	6
7	8	9



1		2		3				
		4		5		6		
				7		8		9

# Inception V3 model

- Deep learning model based on Convolutional Neural Networks
- Used for image classification
- Released in year 2015
- It has 42 layers


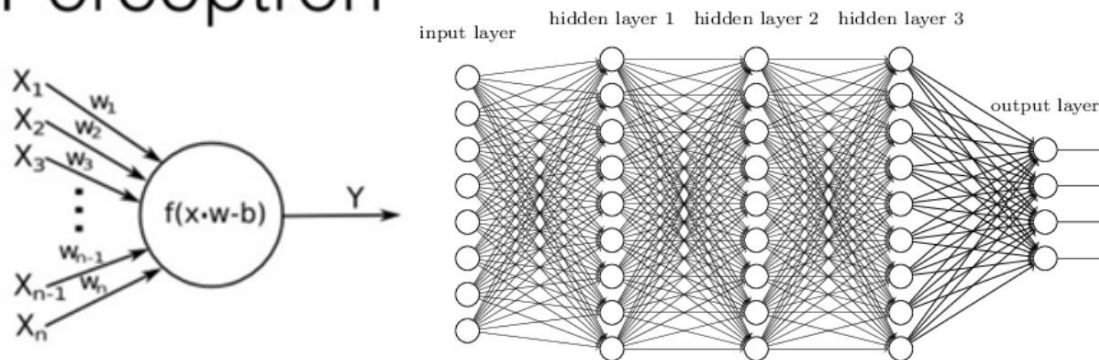


# Overview of the notebook

# Tutorial IV (1)

- 1) Load necessary **libraries** (common libraries and personal modules)
- 2) Images
- 3) **Convolutions**

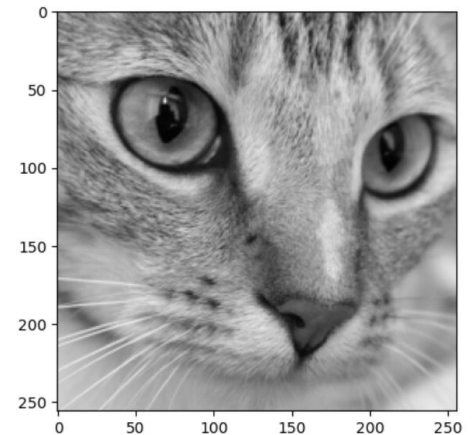
## Perceptron



# Tutorial IV (2)

- Pre-processing of the image

- Gray-scale, cropping, float conversion, normalization
- Add dimensions (batch, channel, height, width) (in `get_convolved`)
- Convert Numpy to pytorch (in `get_convolved`)



- Define the convolution

- Forward model (`class Model`)
- Apply 4 convolutions one after each other (inside `class Model`, call `conv_2d` function)

- Define the filter

- Identity filter
- Convert to `np.array` (in `get_convolved`)
- Add dimensions
- Convert Numpy to pytorch

```
flt_mtx = [  
    [ 0, 0, 0, 0, 0,],  
    [ 0, 0, 0, 0, 0,],  
    [ 0, 0, 1, 0, 0,],  
    [ 0, 0, 0, 0, 0,],  
    [ 0, 0, 0, 0, 0,],  
] # identity transformation
```



# Tutorial IV (3)

- Use it ! (`ims_convolved = get_convolved(img_raw, flt_mtx)`)
  - You get back 5 figures

- Exercise :

1. experiment with different filters and understand what they do, e.g.:

- identity transformation
- identity transformation with positive non-unit values
- identity transformation with negative unit value
- identity transformation off center
- blurring with box filter
- edge detection with + and - bands
- try whatever you like

2. experiment with convolution parameters:

- padding = 1, 2, 3
- stride = 2
- dilation = 2

0	0	0	0	0
0	0	0	0	0
0	0	30	0	0
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
0	0	0	0	0
0	0	-1	0	0
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0

# Tutorial IV (4)

filter type	effect
gaussian	blurring
first derivative of gaussian	detection of edges
second derivative of gaussian	detection of peaks

- **Most common filters**

- Define 1D functions
- Create 2D filters by repeating the 1D filters along axis 0 (`np.tile`)
- Multiply by `transpose()` to get the horizontal dimension (filter size does not change)
- Use them ! (`ims_convolved = get_convolved(img_raw,flt_mtx)`)

## 4) Homework (leave it for now)

# Tutorial IV (5)

5) Load a pretrained model (inception V3) from torchhub

6) Test the model

- Preprocessing of the image (cropping, shuffle sizes, totensor, normalize adds batch size) → [1,3,299,299]
- Deactivate the gradient (eval mode)
- Get the logits (1000 values), then the probabilities (applying softmax)
- Print out the 5 most probable classes
- Do the same with 100 classes

# Tutorial V : Transfer Learning

# Introduction

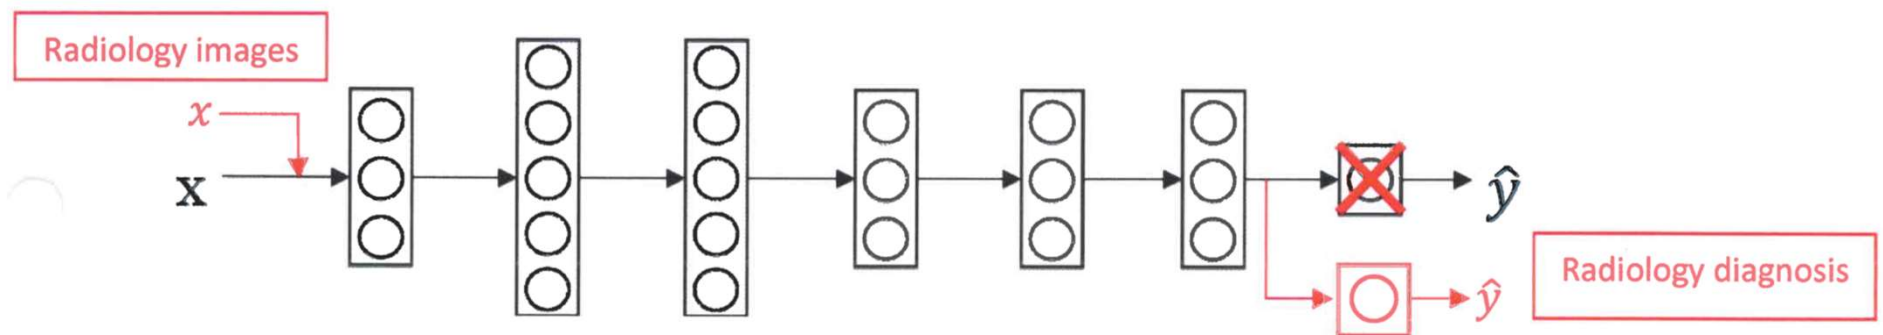
- **Goal** : use the Inception model to classify images of different nature, learn how to save a model
- **Program** : inverted classroom style
  - Theory
  - Overview to get the big picture of the notebook
  - Work alone ()
  - Work in groups of 4 (45 minutes, then 2 minute presentation)
- **Technical** : Google Colab, Pytorch



# Theory

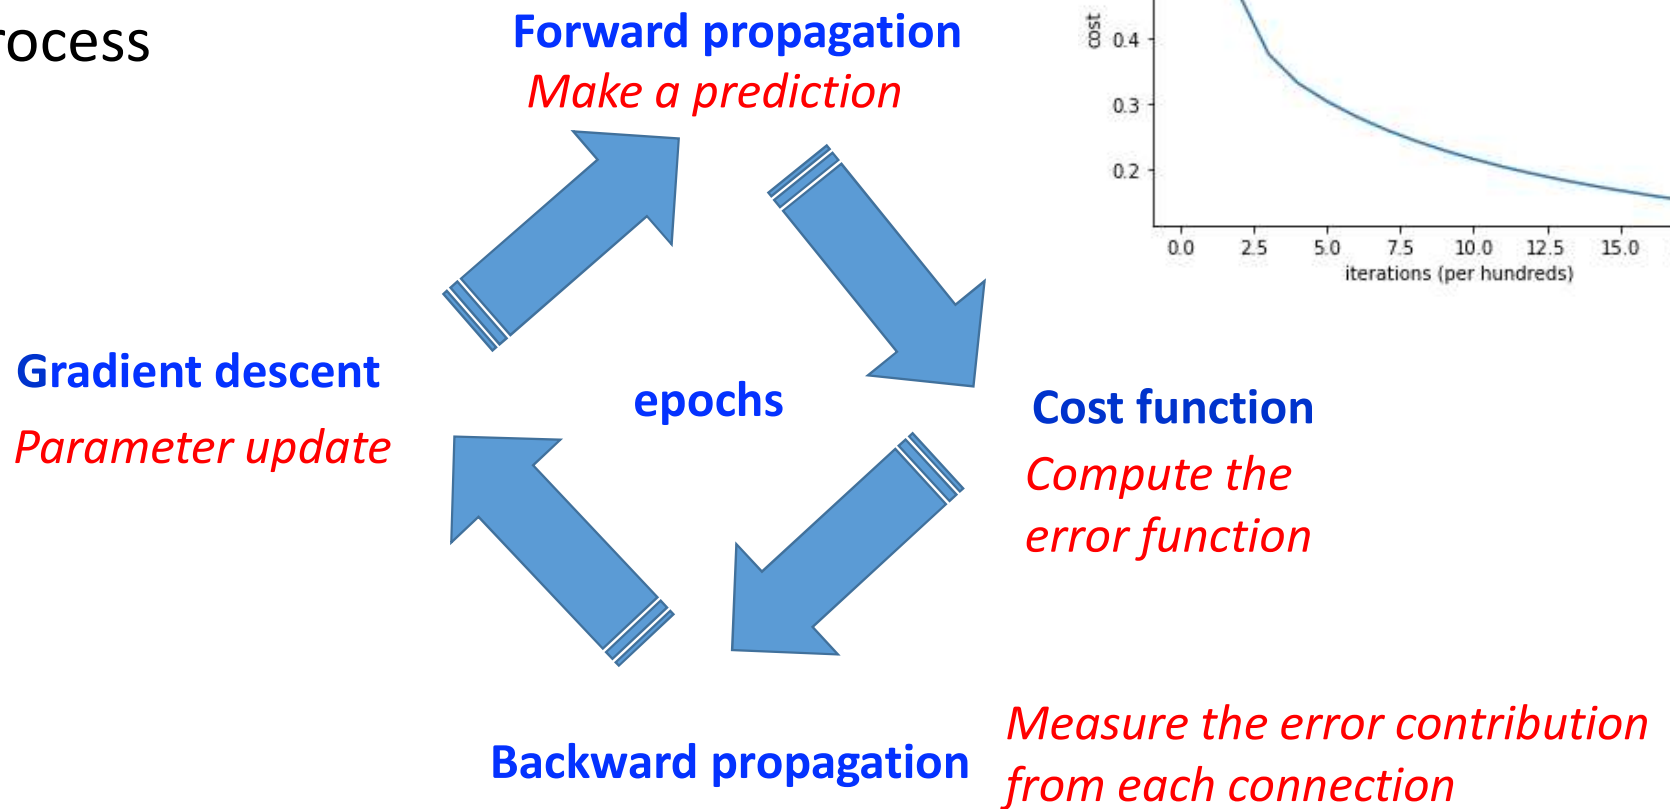
# Transfer Learning

- Try to find an existing neural network that accomplishes a **similar task** to the one you are trying to tackle
- reuse the lower layers of this network
  - Output layer should usually be replaced
- **Speeds up** training and requires much fewer training data

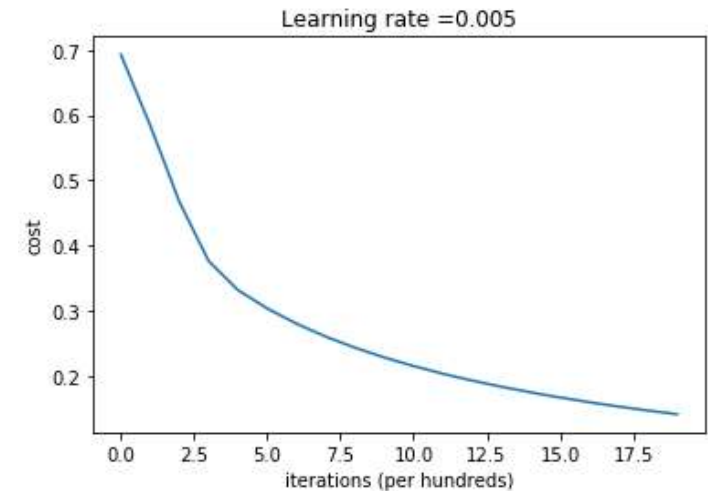


# Training Loop

- *Iterative* process

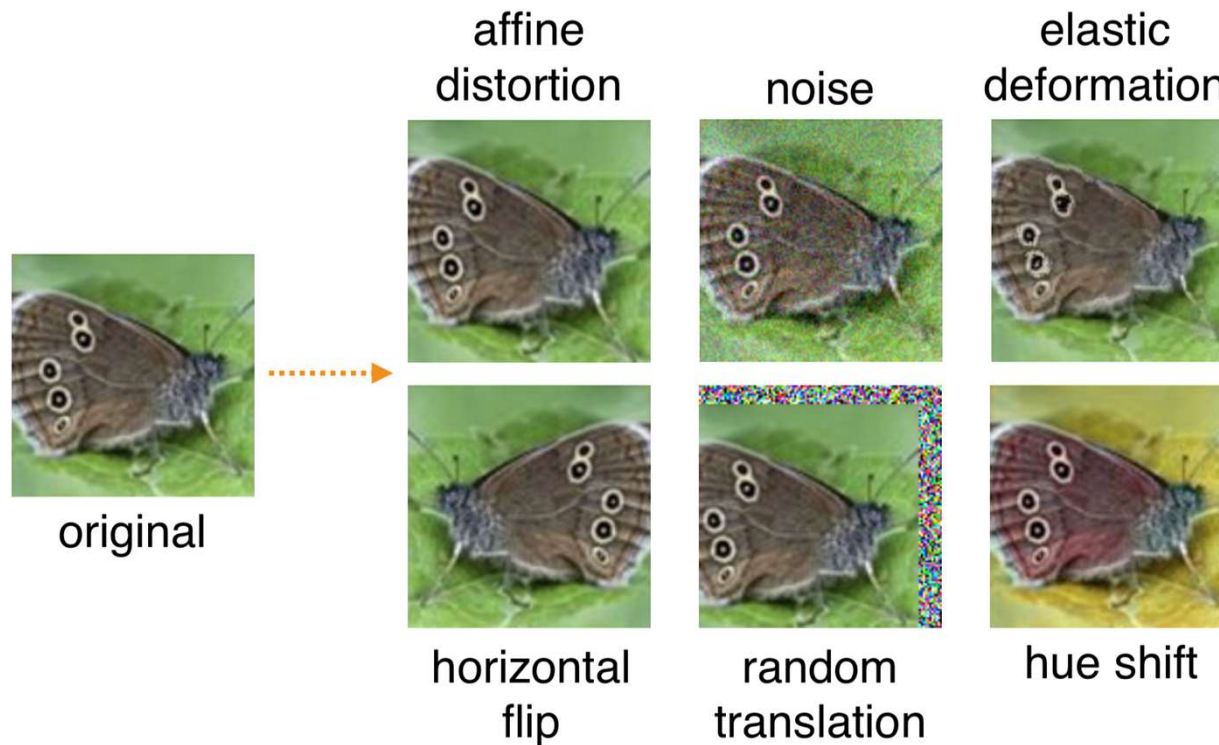


Learning curve



# Dataset Augmentation

- Apply **realistic transformations** to data to create new synthetic samples, with same label



# Overview of the notebook



# Tutorial V (1)

1) Load necessary **libraries** (common libraries and personal modules)

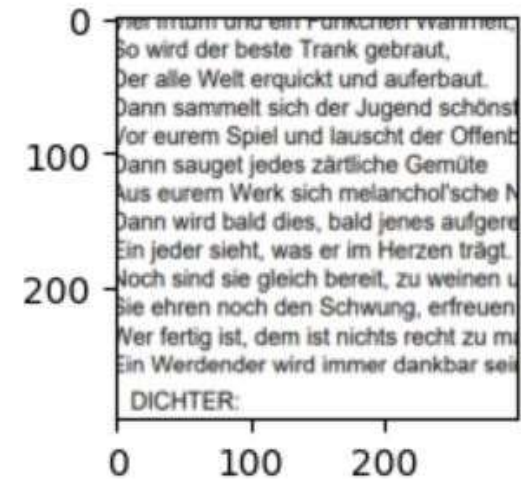
## 2) Transfer Learning

- Load the **inception model** (**base\_model**)
- Build **new model** using the **base\_model** (**model**)
  - Define a head function (in\_features, n\_classes=2) : use of sigmoid
- **Optimization** : define the loss function (**criterion**) and the optimizer
- **Helper functions** to get the prediction (**get\_predictions**) and to compute the batch accuracy (**calculate\_accuracy\_batch**)

# Tutorial V (2)

## 3) Dataset

- **Images** :
  - Get them from **ML3** folder
  - Preprocessing using **transforms.Compose (resize, tensor, normalize)**
  - Shape **[3,299,299]**
  - Transform to numpy array for display (**im\_numpy**)
- **Labels**
- **Split** dataset into train/test samples
  - **Torch.utils.data.random\_split**
  - **Train\_test\_split** with **stratify** enabled from scikit-learn
- Create **data loaders** for training/val (beware shuffle param)
- Example : batch of 10 images, plot them, print logits and output classes (**res**)



# Tutorial V (3)

## 4) Training

- **Train\_model** function
  - Contains the loop on **epochs**
  - Calls the **train** and **validate** functions (beware the params)
  - Save **history** (loss, accuracy) and model for a given epoch
- **Train function** : reset gradients, compute logits and loss, compute gradients (backward prop), update parameters, calculate accuracy of the batch (helper functions), returns train loss and train accuracy
- **Validate function**: structure BUT (no optimizer, no backward, no update of the params), returns test loss and test accuracy
- **Plot\_history** function
- Run it ! **history = train\_model(...)** using 70 epochs

# Tutorial V (4)

## 5) Load trained variables from checkpoint

- Choose epochs values
- Load corresponding models
- Call validate function to get the validation loss and validation accuracy (see how it evolves)

## 6) Save final model for inference

## 7) Inference :

- Load the model, eval() mode
- Get an image, preprocess (convert to tensor, add batch dimension)
- Get the logits and associated class

# Tutorial V (4)

## 8) Improve the results : data augmentation

- Load images from **ML3** folder
- Preprocess (**resize**, **Randomcrop**, **tensor**, **normalize**)
- Convert to Numpy for plotting purposes (**im\_numpy**)
- The rest of the code is similar to previous code

## 9) Exercise (groups of 4) (35 min), 2min presentation

# Tutorial VI : Transfer Learning

# Introduction

- **Goal** : use Recurrent Neural Networks to predict and generate text sequence
- **Program** : inverted classroom style
  - Theory
  - Overview to get the big picture of the notebook
  - Work alone
  - Work in groups of 4
- **Technical** : Google Colab, Pytorch



# Theory

# Overview of the notebook

# Tutorial VI (1)

1) Load necessary **libraries** (common libraries and personal modules)

## 2) Text data

- Read the data (**rnn.txt**), print the first 100 words

## 3) Build the **dataset**

- 2 dictionaries : word → id (**dictionary**) and id → word (**reverse\_dictionary**)
  - Build\_dictionaries function
- Vocabulary size = 493 (0=most common word)
- Helper functions to get sequences of int or words (**text\_to\_ints**, **ints\_to\_text**)
- Print example : first 100 words (or int), length of input data=2118 (**words\_as\_int**)

# Tutorial VI (2)

## 3) Data streaming

- Create dataset using **WordDataSet** class to create blocks of text
  - Block length =  $n\_input + 1 = 3 + 1 = 4$
- Create **DataLoader** with **batch size=50**, **preprocess** data (separate input and target sequences, **stack** data and convert Numpy → **tensors**, put them to **GPU**)
  - Length of dataset = Number of blocks in sequence =  $\text{Total length} / \text{Block length} = 2118 / 4 = 529$
- Example : print the 50 samples that are in the 1st batch

# Tutorial VI (3)

## 4) Construct model

- Create **class RNN**
  - Embedding layer (vocab\_size, embedding\_dim=128)
  - Loop to add the 3 **LSTM** layers
  - **FC** layer with vocab\_size
- Define sequence of 3 words (**n\_input**), define dimension of 3 LSTM, call the RNN class
- Investigate the model using **Tensorboard**
  - Create an input of size=5 (**x**), transform to tensor, add batch
  - **SummaryWriter** to save the model and be able to open it with tensorboard
  - Print the output size of **y** : [5,1,493]
- Test the **NOT trained model** and see that it is bad
  - Get the first batch (break), apply the model, get the predictions, compare to true

# Tutorial VI (4)

## 5) Train the model

- Params : `n_input = 3`, `batch_size=50` , one LSTM layer (128), `n_epochs=200`
- Create the `data loader` (preprocess data)
- `Optimization` part : `criterion`, `optimizer` (RMSprop, why?)
- `Training loop` on epochs, then on batches
  - Initialize gradients
  - Get the output (`seq_len`, `batch_size`, `vocab_size`), reshape it to (`seq_len*batch_size`, `vocab_size`)
  - Reshape labels (`seq_len`, `batch_size`) to (`seq_len*batch_size`)
  - Compute the loss between output and true labels
  - Backward prop, param update
  - Compute loss and accuracy
- `Plot` loss and accuracy

# Tutorial VI (5)

## 6) Generate text with RNN

- Function to generate text (`gen_long`)
  - Input parameters : model, input sequence, number of words to generate (128)
  - `No_grad()` because we are in an evaluation mode
  - Loop on number of words, convert to tensor, predict, ...)



Back-up slides

- [https://www.tensorflow.org/api\\_docs/python/tf/keras/applications](https://www.tensorflow.org/api_docs/python/tf/keras/applications)